

---

# Chapter 4

## 메모리 관리

18 January, 2021

Minguk Choi

koreachoi96@gmail.com

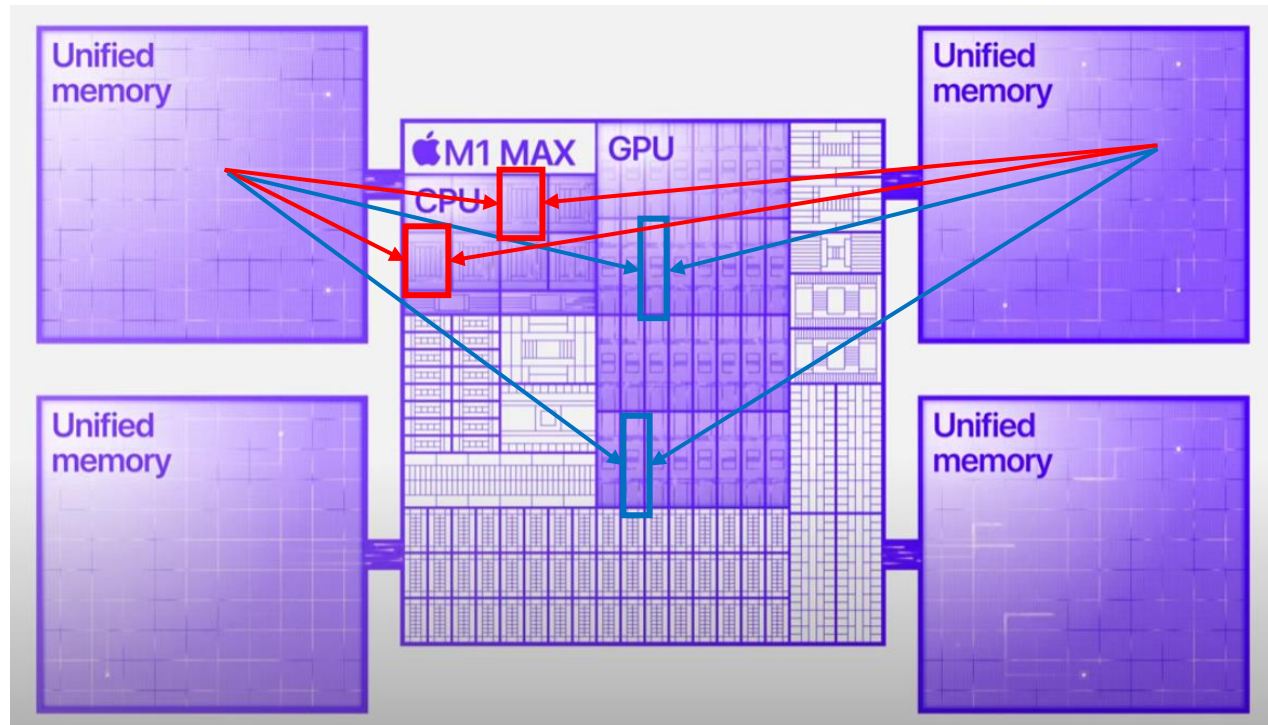
# Contents

---

- Physical Memory Data Structure
  - UMA/NUMA
  - Node, Zone, Page
- Buddy Allocator
  - alloc\_pages / free\_pages()
  - vmalloc() / vfree()
  - practice1: alloc\_pages() vs vmalloc()
- Slab Allocator
  - kmalloc() / kfree()
  - practice2: kmem\_cache\_alloc() vs kmalloc()
- Kernel Stack
  - practice3: kernel stack overflow


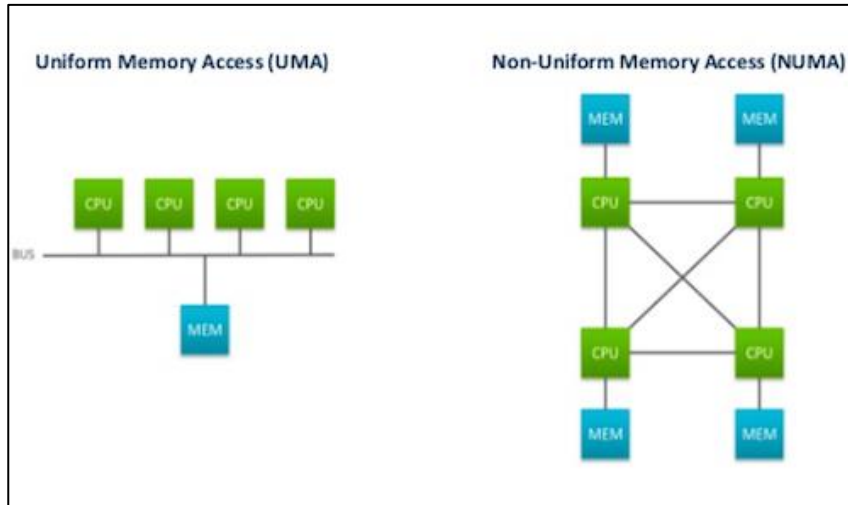
# Physical Memory Data Structure

- Apple M1 Max SoC



# Physical Memory Data Structure

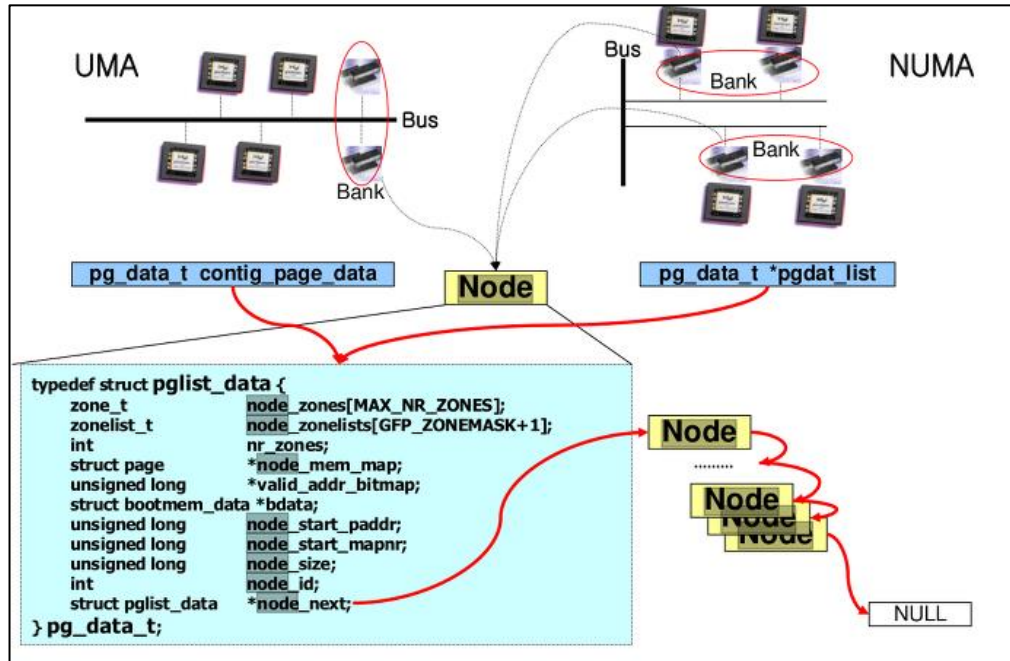
## ■ SMP/NUMA



- SMP(UMA)
  - ✓ Symmetric Multiprocessing
  - ✓ 모든 CPU와 메모리가 하나의 입출력 버스 등을 공유
  - ✓ 병목현상 발생
- NUMA
  - ✓ Non-Uniform Memory Access
  - ✓ CPU를 그룹으로 나누고, 각 그룹에 별도로 지역 메모리 할당
  - ✓ CPU가 어떤 메모리를 접근하느냐에 따라 성능 차이 발생

# Physical Memory Data Structure

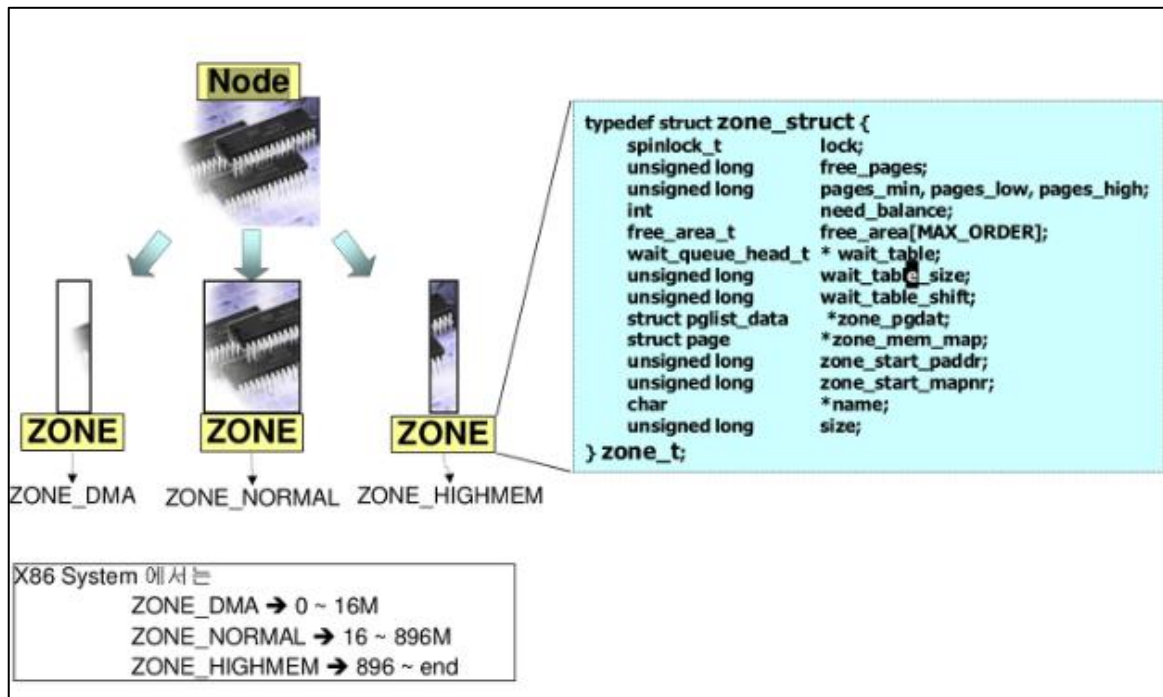
- Node



- Bank
  - ✓ 리눅스에서 접근 속도가 같은 메모리의 집합
  - ✓ UMA: single Bank
  - ✓ NUMA: non-single Banks
- Node
  - ✓ 리눅스에서 Bank를 표현하는 자료구조

# Physical Memory Data Structure

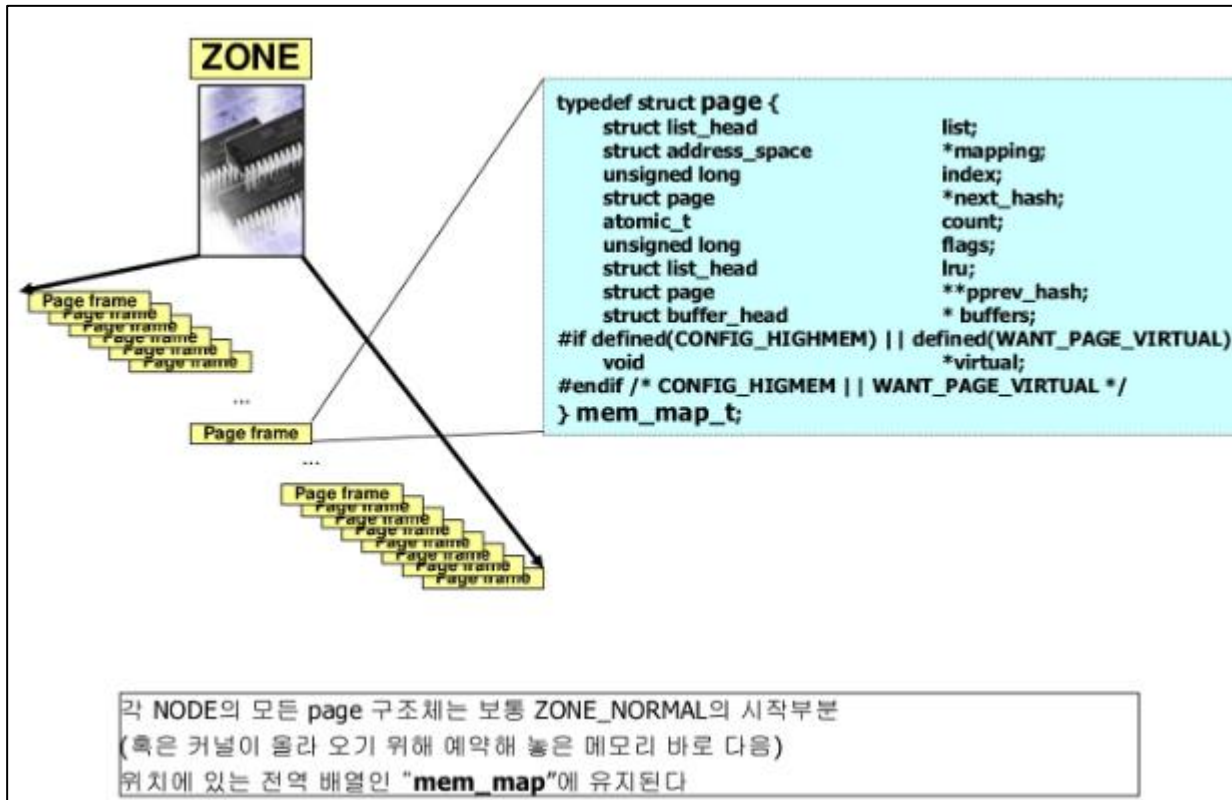
## ■ Zone



- 커널이 모든 메모리를 동일하게 취급?
  - ✓ 일부 하드웨어의 한계로 불가
    - 특정 메모리 주소로만 DMA 수행가능
    - 일부 메모리는 커널 주소 공간에 상주 불가
- ZONE
  - ✓ ZONE\_DMA: DMA 수행가능
  - ✓ ZONE\_NORMAL
  - ✓ ZONE\_HIGHMEM: 커널 주소공간 상주X  
페이지 동적으로 연결
- 32bit-x86
  - ✓ ZONE\_DMA: 0MB ~ 16MB
  - ✓ ZONE\_NORMAL: 16MB ~ 896MB
  - ✓ ZONE\_HIGHMEM: 896MB ~

# Physical Memory Data Structure

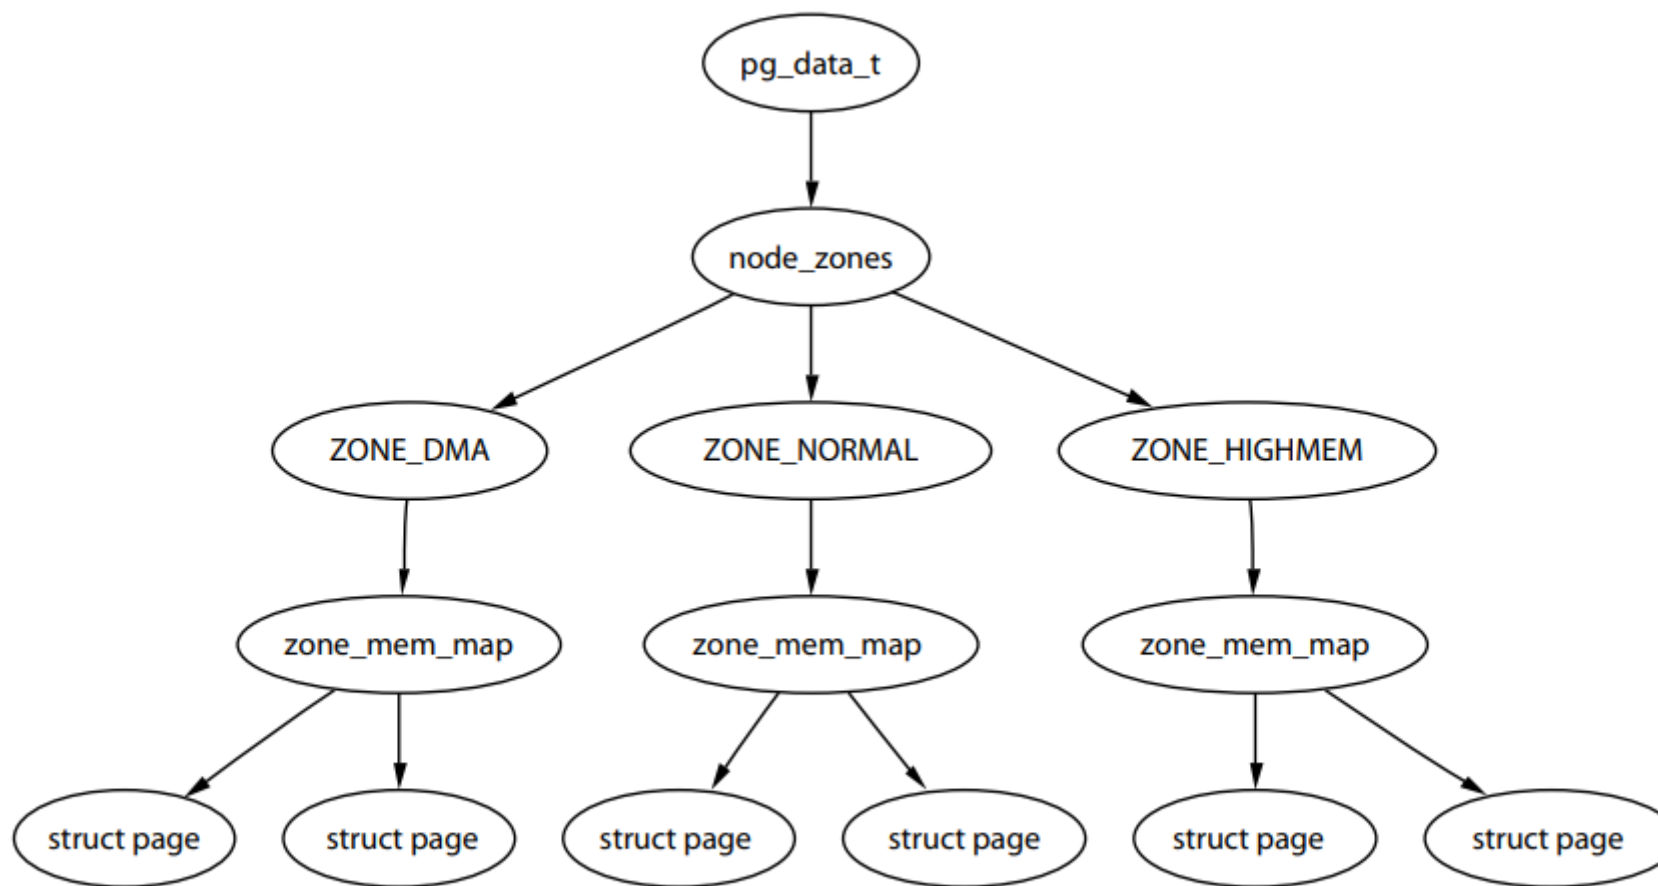
## ■ Page



- 물리적 메모리 최소 관리단위
  - ✓ 커널: Page
  - ✓ Why? MMU(Memory Management Unit)
    - 페이지 단위로 가상 메모리 주소를 물리적 메모리 주소로 변환
- Page Size
  - ✓ 32bit: 4KB
  - ✓ 64bit: 8KB
- Struct page
  - ✓ 가상 페이지가 아닌, 물리적 페이지 표현
    - Struct page의 내용은 일시적 (ex. swap)

# Physical Memory Data Structure

---



**Figure 2.1.** Relationship Between Nodes, Zones and Pages



# Buddy Allocator

- `alloc_pages / free_pages()`
  - contiguous VA, contiguous PA

```
struct page * alloc_pages(unsigned int gfp_mask, unsigned int order)
```

Allocates  $2^{\text{order}}$  number of pages and returns a struct page.

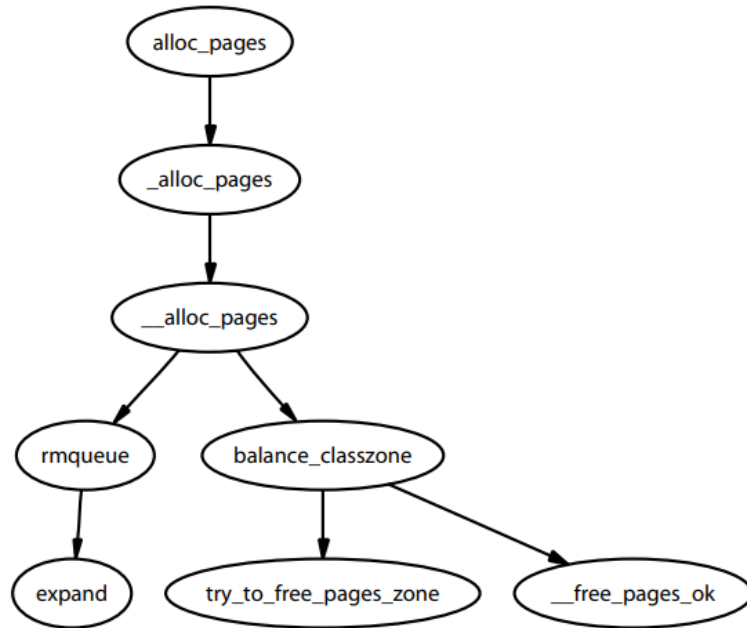


Figure 6.2. Call Graph: `alloc_pages()`

```
void __free_pages(struct page *page, unsigned int order)
```

Frees an order number of pages from the given page.


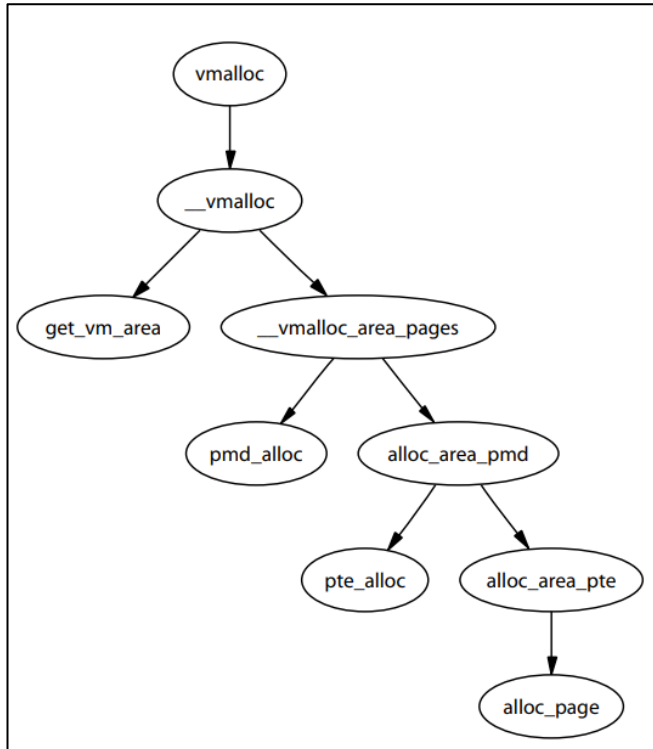


Figure 6.4. Call Graph: `__free_pages()`

# Buddy Allocator

- `vmalloc()` / `vfree()`
  - contiguous VA, but do not guarantee contiguous PA

```
void * vmalloc(unsigned long size)  
    Allocates a number of pages in vmalloc space that satisfy the requested size.
```



```
void vfree(void *addr)  
    Frees a region of memory allocated with vmalloc(), vmalloc_dma() or  
    vmalloc_32()
```

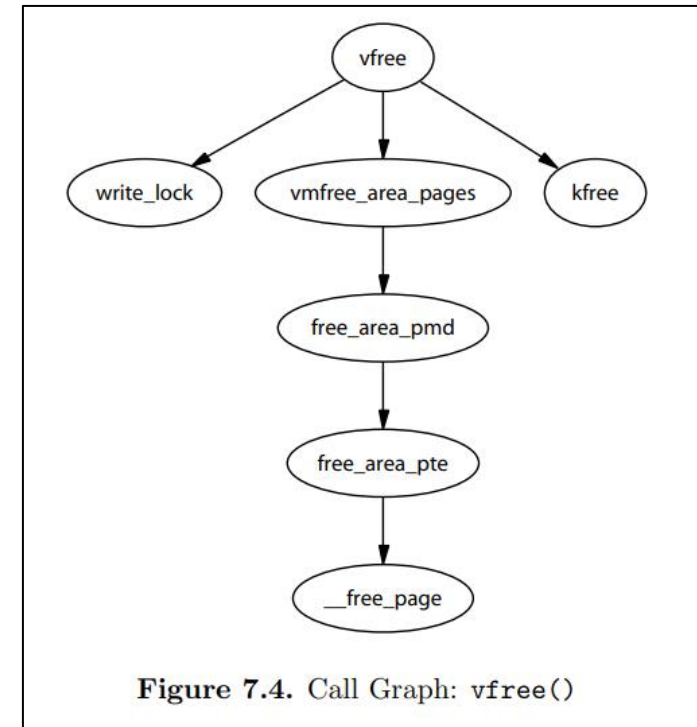
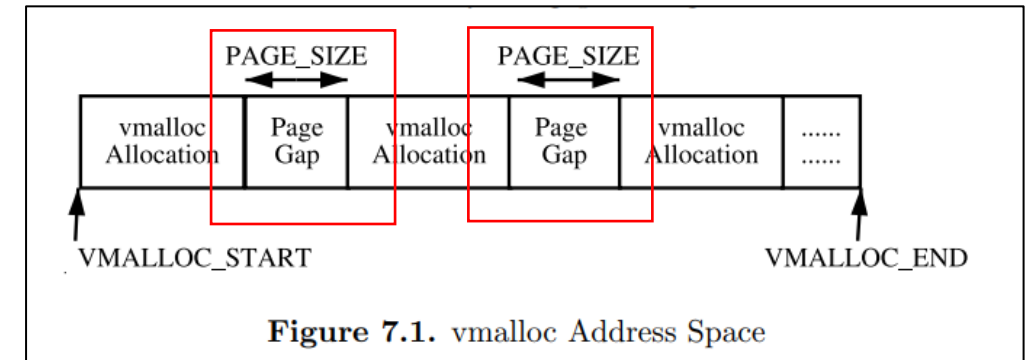
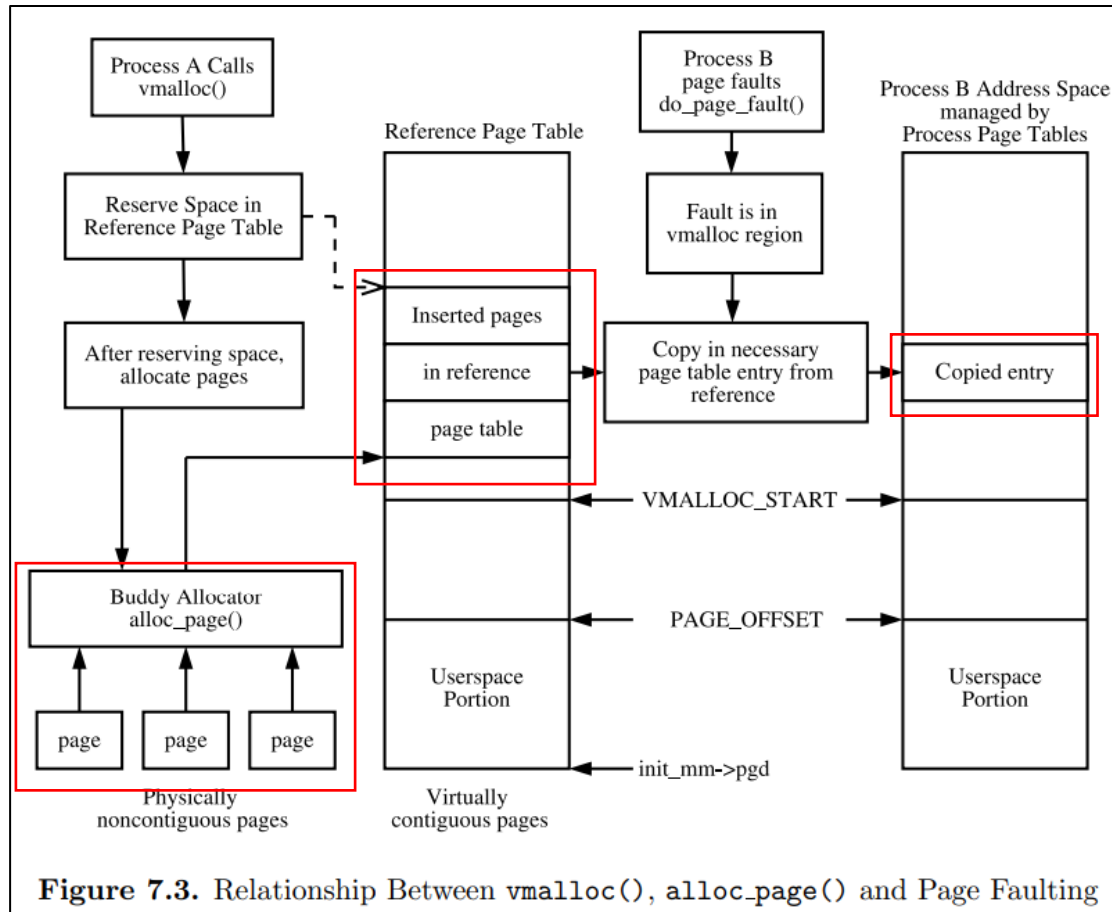


Figure 7.4. Call Graph: `vfree()`

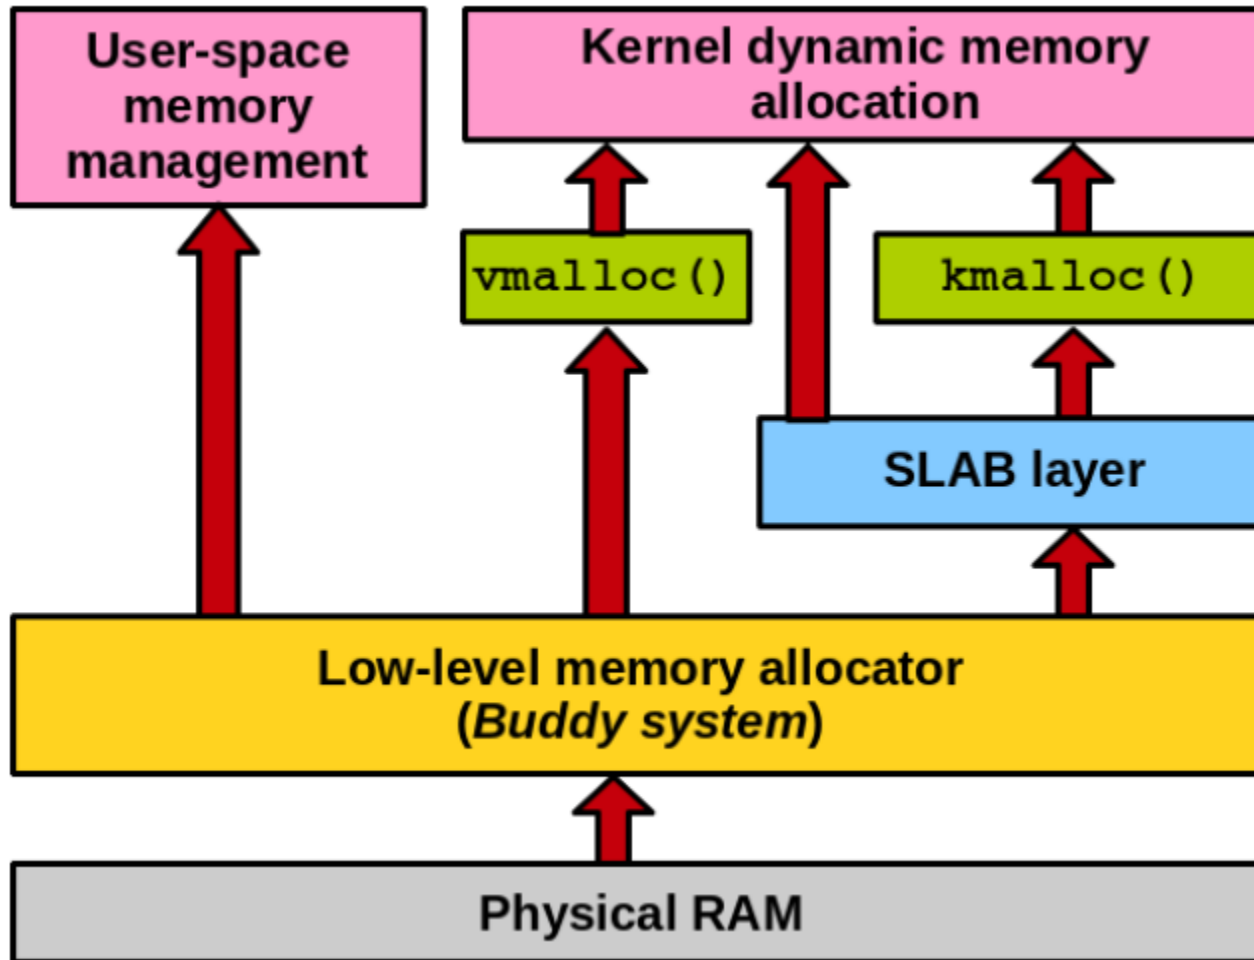
# Buddy Allocator

## ■ vmalloc()



# Buddy Allocator

---



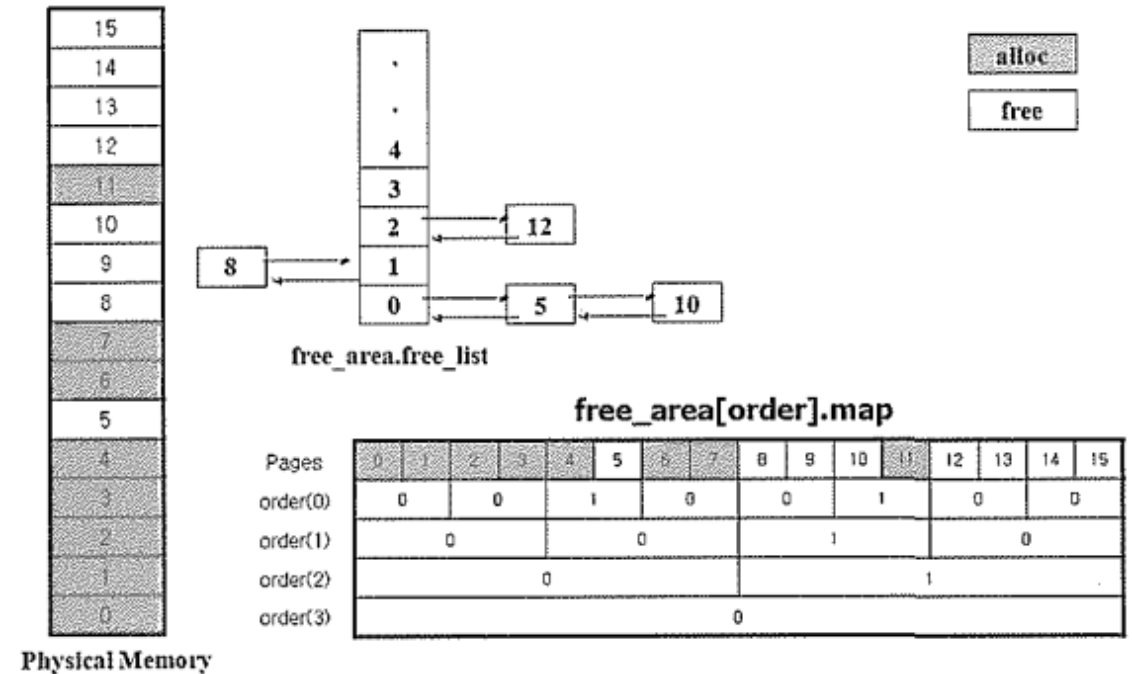
# Buddy Allocator

- Buddy Allocator mechanism

```
/* ~/include/linux/mmzone.h
#define MAX_ORDER 10

struct zone {
    ...
    struct free_area free_area[MAX_ORDER];
    ...
};


struct free_area {
    struct list_head free_list;
    unsigned long *map;
};
```



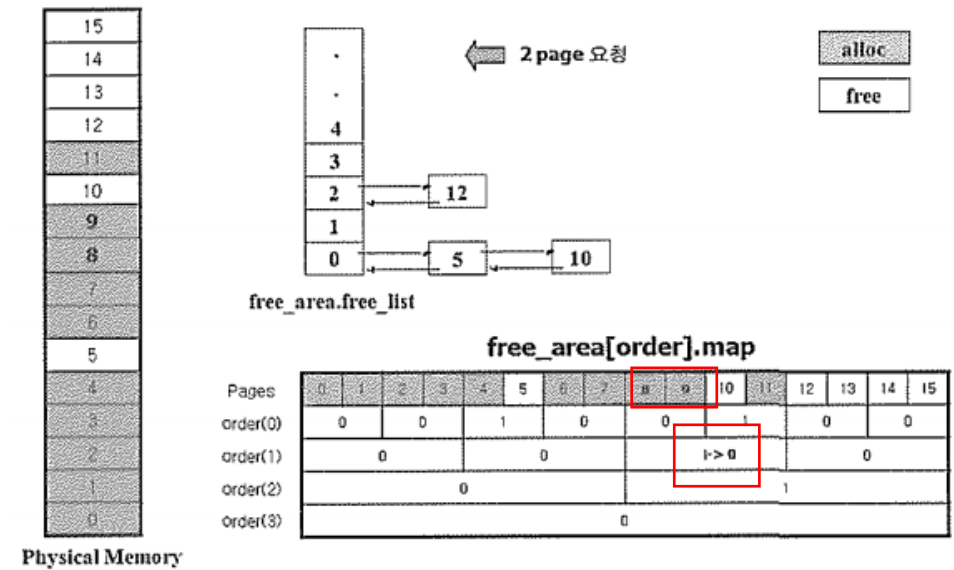
(1)

# Buddy Allocator

- Buddy Allocator mechanism
  - On 2 pages are requested




(1)



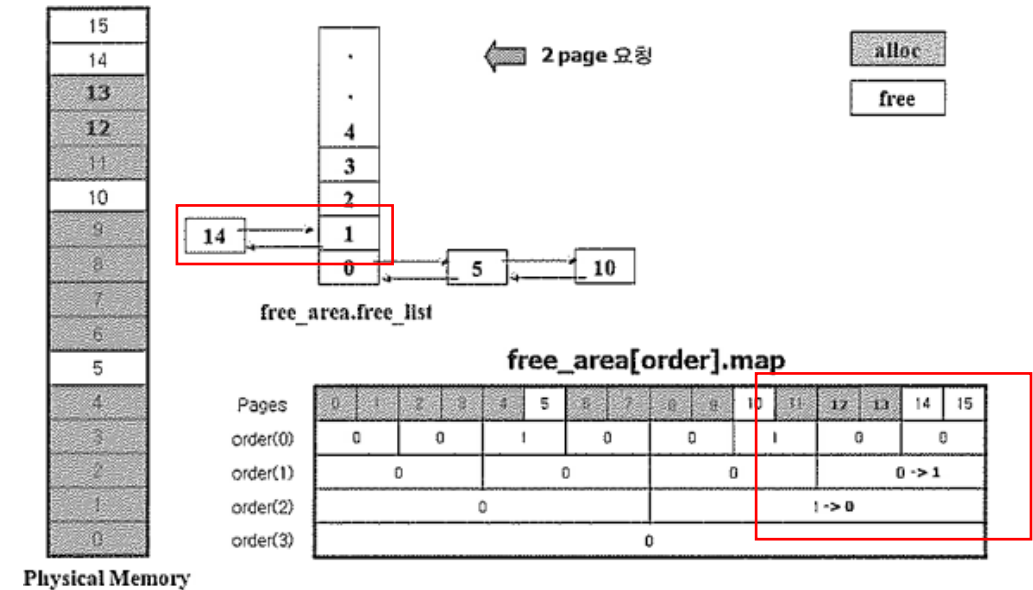
(2)

# Buddy Allocator

- Buddy Allocator mechanism
  - On another 2 pages are requested



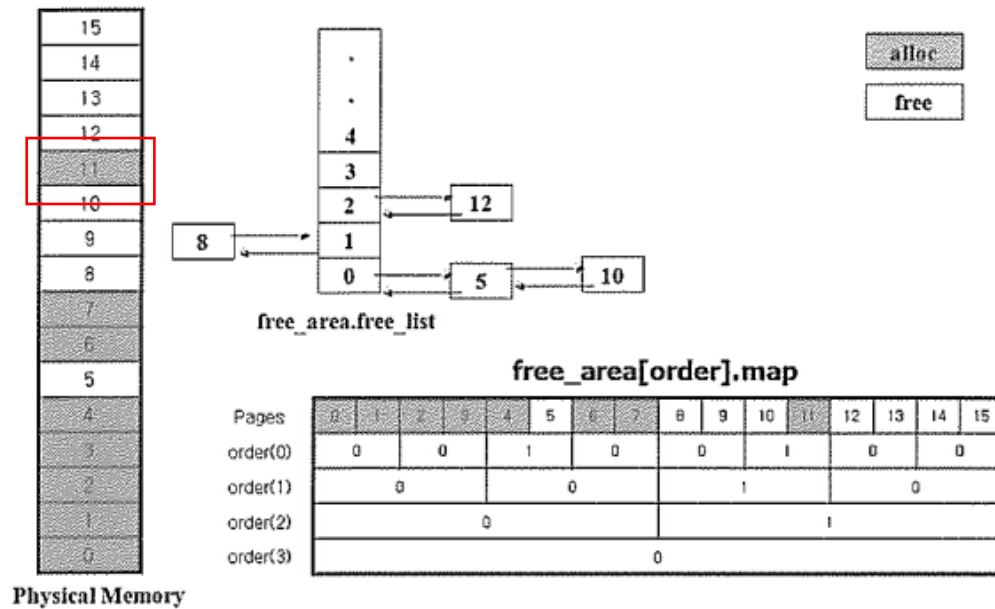
(1)




(2)

# Buddy Allocator

- Buddy Allocator mechanism
  - On page 11 are freed



(1)




(2)



# Buddy Allocator

- Lazy Buddy allocator
  - Zone Watermarks



```
37 typedef struct zone_struct {
41     spinlock_t      lock;
42     unsigned long    free_pages;
43     unsigned long    pages_min, pages_low, pages_high;
44     int              need_balance;
```

**pages\_high** After **kswapd** has been woken to start freeing pages, it will not consider the zone to be “balanced” when **pages\_high** pages are free. After the watermark has been reached, **kswapd** will go back to sleep. In Solaris, this is called **lotsfree**, and, in BSD, it is called **free.target**. The default for **pages\_high** is three times the value of **pages\_min**.

**pages\_low** When the **pages\_low** number of free pages is reached, **kswapd** is woken up by the buddy allocator to start freeing pages. This is equivalent to when **lotsfree** is reached in Solaris and **freemin** in FreeBSD. The value is twice the value of **pages\_min** by default.

**pages\_min** When **pages\_min** is reached, the allocator will do the **kswapd** work in a synchronous fashion, sometimes referred to as the *direct-reclaim* path. Solaris does not have a real equivalent, but the closest is the **desfree** or **minfree**, which determine how often the pageout scanner is woken up.

# Buddy Allocator

## ■ alloc\_pages vs vmalloc()

```
static int __init kmalloc_init(void) {  
  
    pg_mem = alloc_pages(GFP_KERNEL, 2);  
    vm_mem = vmalloc(32768);  
  
    print_bulk_address(vm_mem, "vmalloc");  
    print_bulk_address(page_address(pg_mem), "alloc_pages");  
    printk("");  
  
    return 0;  
}
```

```
void print_bulk_address(void* mem, char* name){  
    int i, idx=0;  
    printk("");  
  
    for (i=0; i<8; i++){  
        idx = i*4096;  
        printk("%12s %5d: [virtual] %12px [page] %12llx [physical] %12llx",  
               name, idx, mem+idx, virt_to_page(mem+idx), virt_to_phys(mem+idx));  
    }  
}
```

```
static void __exit kmalloc_exit(void) {  
    vfree(vm_mem);  
    free_pages(page_address(pg_mem), 2);  
}
```

```
mingu@mingu-VirtualBox:~/module_vmalloc$ sudo dmesg -c  
  
[ 709.920936]      vmalloc      0: [virtual] fffffb19b41d21000 [page] fffffe63e87074840 [physical] 2061c1d21000  
[ 709.920938]      vmalloc    4096: [virtual] fffffb19b41d22000 [page] fffffe63e87074880 [physical] 2061c1d22000  
[ 709.920939]      vmalloc    8192: [virtual] fffffb19b41d23000 [page] fffffe63e870748c0 [physical] 2061c1d23000  
[ 709.920940]      vmalloc   12288: [virtual] fffffb19b41d24000 [page] fffffe63e87074900 [physical] 2061c1d24000  
[ 709.920941]      vmalloc   16384: [virtual] fffffb19b41d25000 [page] fffffe63e87074940 [physical] 2061c1d25000  
[ 709.920942]      vmalloc   20480: [virtual] fffffb19b41d26000 [page] fffffe63e87074980 [physical] 2061c1d26000  
[ 709.920943]      vmalloc   24576: [virtual] fffffb19b41d27000 [page] fffffe63e870749c0 [physical] 2061c1d27000  
[ 709.920944]      vmalloc   28672: [virtual] fffffb19b41d28000 [page] fffffe63e87074a00 [physical] 2061c1d28000  
  
[ 709.920946]  alloc_pages      0: [virtual] fffff913a8c15c000 [page] fffffe5bd04305700 [physical] 10c15c000  
[ 709.920947]  alloc_pages    4096: [virtual] fffff913a8c15d000 [page] fffffe5bd04305740 [physical] 10c15d000  
[ 709.920948]  alloc_pages    8192: [virtual] fffff913a8c15e000 [page] fffffe5bd04305780 [physical] 10c15e000  
[ 709.920949]  alloc_pages   12288: [virtual] fffff913a8c15f000 [page] fffffe5bd043057c0 [physical] 10c15f000  
[ 709.920950]  alloc_pages   16384: [virtual] fffff913a8c160000 [page] fffffe5bd04305800 [physical] 10c160000  
[ 709.920951]  alloc_pages   20480: [virtual] fffff913a8c161000 [page] fffffe5bd04305840 [physical] 10c161000  
[ 709.920952]  alloc_pages   24576: [virtual] fffff913a8c162000 [page] fffffe5bd04305880 [physical] 10c162000  
[ 709.920953]  alloc_pages   28672: [virtual] fffff913a8c163000 [page] fffffe5bd043058c0 [physical] 10c163000
```

- alloc\_pages()
  - contiguous VA, contiguous PA
- vmalloc()
  - contiguous VA, **contiguous PA**

# Buddy Allocator

---

- **RACE CONDITION:** Alloc\_pages vs vmalloc()

While(1)

vmalloc(random()/32\*PAGE\_SIZE)  
vfree()

alloc\_pages(GFP\_KERNEL, random()%4)  
free\_pages()

malloc(random()/64\*PAGE\_SIZE)  
~~free()~~

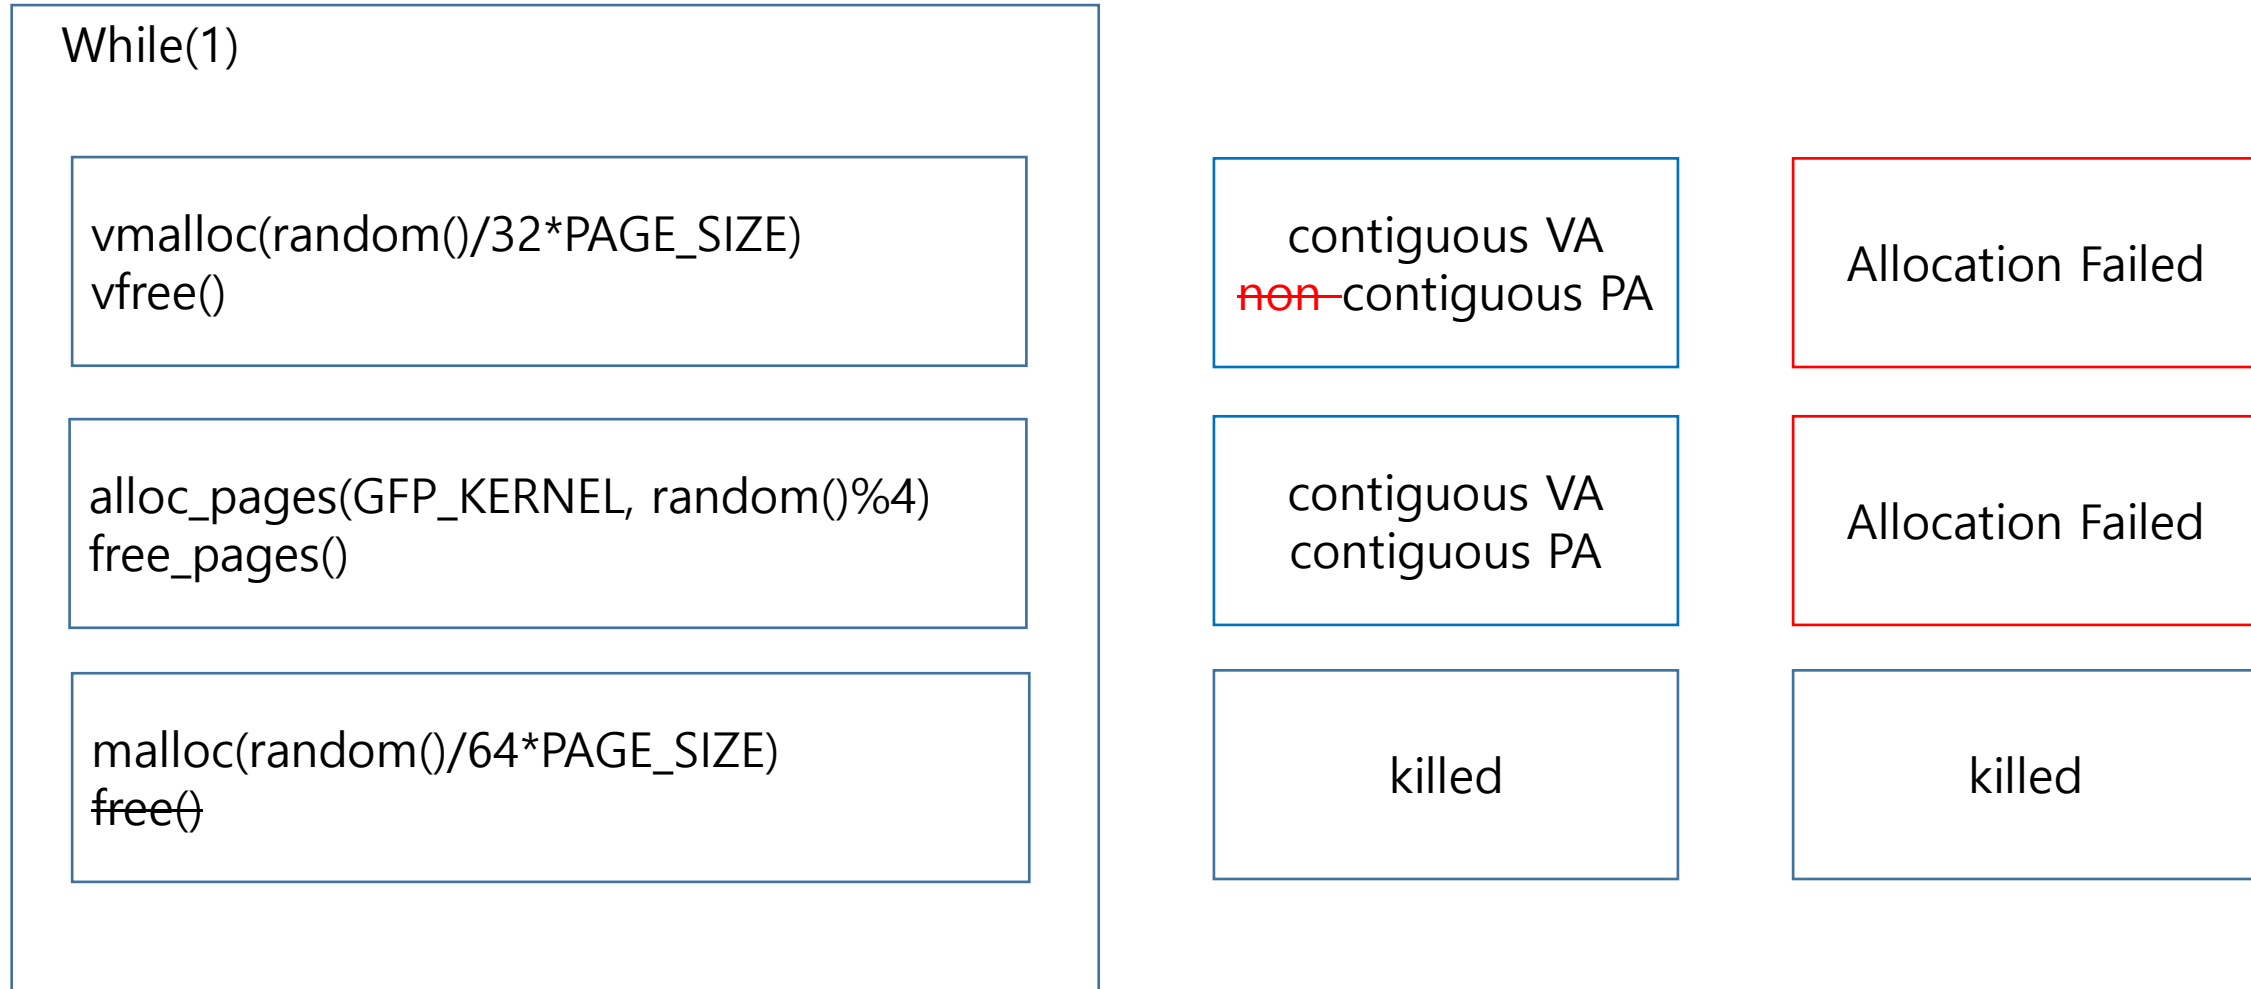
contiguous VA  
non-contiguous PA

contiguous VA  
contiguous PA

killed


# Buddy Allocator

- **RACE CONDITION:** Alloc\_pages vs vmalloc()




# Slab Allocator

- 커널 내부에서 가장 빈번한 작업
  - 자료구조(struct task, struct ...) 할당/해제
  - 이를 vmalloc으로 할당/해제 ?
    - ✓ Memory fragmentation
    - ✓ Allocate & Free overhead
    - ✓ Non-contiguous PA <-> contiguous VA
    - ✓ Need contiguous PA for hardware device
      - 가상 주소 처리하는 메모리 장치 x




- 대신 Free list를 사용
    - 자료구조 사용 이후, 해제하지 않고 free\_list에 추가
    - 일종의 객체 Cache
    - But...
      - ✓ 메모리 부족 시, 어떤 캐시 크기를 줄일 것일지 모름
- > Slab Allocator 사용하여, 캐시 관리



# Slab Allocator

## ■ Slab Allocator

- Cache – Slab – Object
- Slab
  - ✓ full / partial / free
  - ✓ allocate on partial -> free -> new
    - Memory fragmentation mitigate



```
190 struct kmem_cache_s {  
193     struct list_head    slabs_full;  
194     struct list_head    slabs_partial;  
195     struct list_head    slabs_free;  
196     unsigned int         objsize;  
197     unsigned int         flags;  
198     unsigned int         num;  
199     spinlock_t           spinlock;
```

```
typedef struct slab_s {  
    struct list_head    list;  
    unsigned long        colouroff;  
    void                *s_mem;  
    unsigned int         inuse;  
    kmem_bufctl_t        free;  
} slab_t;
```

# Slab Allocator

## ■ kmalloc() / kfree()

- 바이트 단위로 할당/해제
- Contiguous VA, contiguous PA

```
kmem_cache_t * kmem_cache_create(const char *name, size_t size,  
size_t offset, unsigned long flags,  
void (*ctor)(void*, kmem_cache_t *, unsigned long),  
void (*dtor)(void*, kmem_cache_t *, unsigned long))  
Creates a new cache and adds it to the cache chain.
```

## ■ Kmem\_cache

- Kmem\_cache\_create()
- Kmem\_cache\_alloc() <- kmalloc()
- Kmem\_cache\_free() <- kfree()
- Kmem\_cache\_destroy()

```
void * kmem_cache_alloc(kmem_cache_t *cachep, int flags)  
Allocates a single object from the cache and returns it to the caller.
```

```
void * kmalloc(size_t size, int flags)  
Allocates a block of memory from one of the sizes cache.
```

```
void kmem_cache_free(kmem_cache_t *cachep, void *objp)  
Frees an object and returns it to the cache.
```

```
void kfree(const void *objp)  
Frees a block of memory allocated with kmalloc.
```

```
int kmem_cache_destroy(kmem_cache_t * cachep)  
Destroys all objects in all slabs and frees up all associated memory before  
removing the cache from the chain.
```



# Slab Allocator

- `kmem_cache_alloc()` vs `kmalloc()`
- `kmem_cache_free()` vs `kfree()`

```
static int __init kmalloc_init(void) {
    kcache = kmem_cache_create("task_struct", sizeof(struct task_struct),
                              ARCH_MIN_TASKALIGN, SLAB_PANIC | SLAB_TRACE, NULL);

    my_tsk1 = kmem_cache_alloc(kcache, GFP_KERNEL);
    my_tsk2 = kmem_cache_alloc(kcache, GFP_KERNEL);

    kr_tsk1 = kmalloc(sizeof(struct task_struct), GFP_KERNEL);
    kr_tsk2 = kmalloc(sizeof(struct task_struct), GFP_KERNEL);

    print_address(kcache, "kcache");
    print_address(my_tsk1, "my_tsk1");
    print_address(my_tsk2, "my_tsk2");
    printk("");

    print_address(kr_tsk1, "kr_tsk1");
    print_address(kr_tsk2, "kr_tsk2");
    printk("");

    return 0;
}
```

```
static void __exit kmalloc_exit(void) {
    kfree(kr_tsk1);
    kfree(kr_tsk2);

    kmem_cache_free(kcache, my_tsk1);
    kmem_cache_free(kcache, my_tsk2);
    kmem_cache_destroy(kcache);
}
```

- `kmem_cache_alloc()`
  - module
  - close VA & PA
- `kmalloc()`
  - Slab
  - Far VA & PA

```
root@mingu-VirtualBox:/home/mingu/module# dmesg -c
```

[ 2838.892640]	kcache:	[virtual]	ffff8f4c8cb1b400	[page]	ffffdf9dc432c6c0	[physical]	10cb1b400
[ 2838.892644]	my_tsk1:	[virtual]	ffff8f4d68228000	[page]	ffffdf9dc7a08a00	[physical]	1e8228000
[ 2838.892645]	my_tsk2:	[virtual]	ffff8f4d6822ca00	[page]	ffffdf9dc7a08b00	[physical]	1e822ca00
[ 2838.892647]	kr_tsk1:	[virtual]	ffff8f4c850d4000	[page]	ffffdf9dc4143500	[physical]	1050d4000
[ 2838.892648]	kr_tsk2:	[virtual]	ffff8f4c8cbe4000	[page]	ffffdf9dc432f900	[physical]	10cbe4000



# Kernel Stack

## ■ Stack

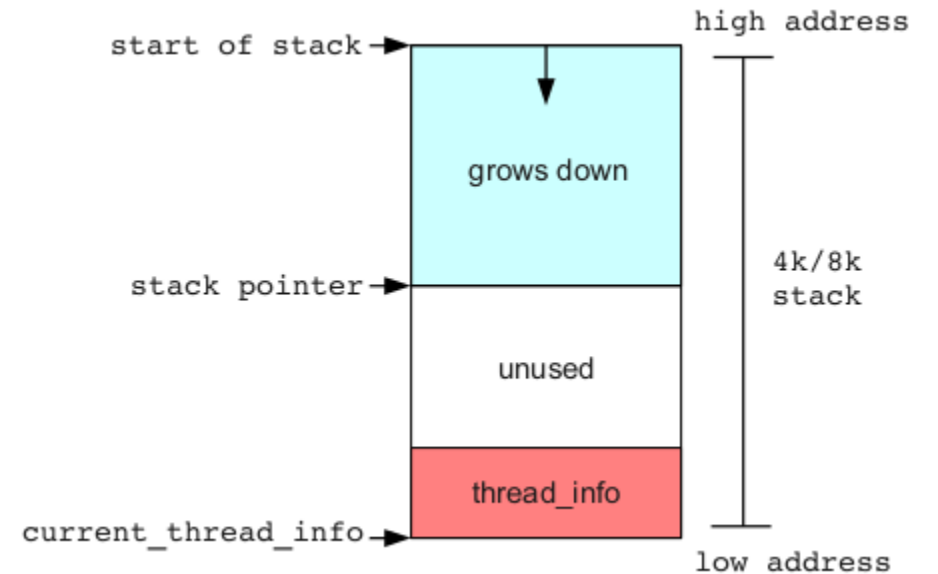
- User Process: big, dynamic
- Kernel: small, fixed
  - ✓ minimize mermory usage
  - ✓ do no not manage kernel stack code

## ■ Kernel Stack size

- $2 * \text{Page size}$ 
  - ✓ 32bit: 8KB, 64bit: 16KB

## ■ Kernel Stack overflow

- therad\_info



# Kernel Stack

## ■ Kernel Stack Overflow

```
static int __init kmallocc_init(void) {  
    int stack_size = 1024; // 1KB  
  
    while(1)  
    {  
        allocate_stack(stack_size);  
        stack_size *= 2;  
    }  
  
    return 0;  
}
```

```
static void allocate_stack(int stack_size){  
    printk("allocate stack, size: %d", stack_size);  
  
    int array[stack_size/4];  
  
    int i;  
    for (i=0; i<array_size; i++)  
        array[i] = 1;  
  
    printk("stack allocated, size: %d", stack_size);  
}
```

```
mingu@mingu-VirtualBox:~/module_stack$ sudo dmesg -c  
[ 5555.374373] allocate stack, size: 1024  
[ 5556.388089] stack allocated, size: 1024  
[ 5556.388099] allocate stack, size: 1024  
[ 5557.411992] stack allocated, size: 1024  
[ 5557.412002] allocate stack, size: 2048  
[ 5558.436207] stack allocated, size: 2048  
mingu@mingu-VirtualBox:~/module_stack$ sudo dmesg -c  
[ 5558.436214] allocate stack, size: 4096  
[ 5559.460679] stack allocated, size: 4096  
[ 5559.460688] allocate stack, size: 8192  
[ 5560.484692] stack allocated, size: 8192  
mingu@mingu-VirtualBox:~/module_stack$ █
```

system doesn't work -> reboot

- Stack Size = 2 \* Page\_Size = 16KB
- Allocate array[16KB]
  - Kernel stack over flow

# Q & A

- Physical Memory Data Structure
  - UMA/NUMA
  - Node, Zone, Page
- Buddy Allocator
  - `alloc_pages` / `free_pages()`
  - `vmalloc()` / `vfree()`
  - practice1: `alloc_pages()` vs `vmalloc()`
- Slab Allocator
  - `kmalloc()` / `kfree()`
  - practice2: `kmem_cache_alloc()` vs `kmalloc()`
- Kernel Stack
  - practice3: kernel size

